

Importing Libraries

```
In [1]: import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, random_split
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt
import itertools
import numpy as np
from torchsummary import summary
from torchvision import models
import torch.nn.functional as F
#import seaborn as sns
#from sklearn.metrics import confusion_matrix
```

Set Device (GPU)

```
In [2]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
device
```

```
Out[2]: device(type='cuda')
```

Data Transformations for Training and Testing

```
In [3]: # Data Transformations
transform_train = transforms.Compose([
    transforms.RandomCrop(32, padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])

transform_test = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])

mean = [0.4914, 0.4822, 0.4465]
std = [0.2023, 0.1994, 0.2010]
classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse']
```

Loading CIFAR-10 Dataset

```
In [4]: # Load CIFAR-10 dataset
trainset_full = torchvision.datasets.CIFAR10(root='./data', train=True, download=True)
testset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True)
```

```

# Display sample images
def show_images_from_classes(dataset, classes, mean=[0.485, 0.456, 0.406], std=[
    fig, axs = plt.subplots(3, 4, figsize=(10, 6))
    axs = axs.ravel()
    displayed_classes = set()
    for i in range(len(dataset)):
        img, label = dataset[i]
        if classes[label] not in displayed_classes:
            displayed_classes.add(classes[label])
            img = img * torch.tensor(std).view(3, 1, 1) + torch.tensor(mean).view(
            npimg = np.clip(img.numpy(), 0, 1)
            axs[len(displayed_classes)-1].imshow(np.transpose(npimg, (1, 2, 0)))
            axs[len(displayed_classes)-1].set_title(classes[label])
            axs[len(displayed_classes)-1].axis('off')

        if len(displayed_classes) == len(classes):
            break

# Turn off any extra subplots
for j in range(len(displayed_classes), len(axs)):
    axs[j].axis('off')


plt.show()

# Display one sample image from each class
show_images_from_classes(trainset_full, classes)

```

Files already downloaded and verified

Files already downloaded and verified



Split Dataset into Training and Validation Sets (80-20 split)

```
In [5]: # Split into Training and Validation Sets (80-20 split)
train_size = int(0.8 * len(trainset_full))
val_size = len(trainset_full) - train_size
trainset, valset = random_split(trainset_full, [train_size, val_size])

trainloader = DataLoader(trainset, batch_size=128, shuffle=True, num_workers=2)
valloader = DataLoader(valset, batch_size=128, shuffle=False, num_workers=2)
testloader = DataLoader(testset, batch_size=100, shuffle=False, num_workers=2)
```

Dataset Summary and Class Display

```
In [28]: def dataset_summary():
print("Dataset Summary:")
print(f"Total training samples: {len(trainset_full)}")
print(f"Training set size: {len(trainset)}")
print(f"Validation set size: {len(valset)}")
print(f"Test set size: {len(testset)}")
print("Classes:", trainset_full.classes)

dataset_summary()
```

Dataset Summary:

Total training samples: 50000

Training set size: 40000

Validation set size: 10000

Test set size: 10000

Classes: ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

Plot Training and Validation Accuracy Curves

```
In [7]: def plot_training_and_validation_accuracy(train_acc, val_acc, model_name):
plt.plot(train_acc, label="Training Accuracy")
plt.plot(val_acc, label="Validation Accuracy")
plt.xlabel("Epoch")
plt.ylabel("Accuracy")
plt.title(f"Training and Validation Accuracy Curve - {model_name}")
plt.legend()
plt.show()
```

Evaluate Model on Test Set

```
In [8]: def evaluate_model(model, dataloader):
model.eval()
correct, total = 0, 0
with torch.no_grad():
    for inputs, labels in dataloader:
        inputs, labels = inputs.to(device), labels.to(device)
        outputs = model(inputs)
        _, predicted = outputs.max(1)
        total += labels.size(0)
```

```

        correct += (predicted == labels).sum().item()
    accuracy = 100 * correct / total
    print(f"Test Accuracy: {accuracy:.2f}%")
    return accuracy

```

Training Function with Validation

```

In [29]: # Training function with validation
def train_model_with_validation(model, criterion, optimizer, num_epochs=10):
    train_acc, val_acc = [], []
    for epoch in range(num_epochs):
        # Training
        model.train()
        correct_train, total_train, running_loss = 0, 0, 0.0
        for inputs, labels in trainloader:
            inputs, labels = inputs.to(device), labels.to(device)
            optimizer.zero_grad()
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            # Accuracy and loss tracking
            _, predicted = outputs.max(1)
            total_train += labels.size(0)
            correct_train += (predicted == labels).sum().item()
            running_loss += loss.item()

        # Calculate training accuracy
        epoch_loss = running_loss / len(trainloader)
        epoch_accuracy = 100 * correct_train / total_train
        train_acc.append(epoch_accuracy)

        # Validation
        model.eval()
        correct_val, total_val = 0, 0
        with torch.no_grad():
            for inputs, labels in valloader:
                inputs, labels = inputs.to(device), labels.to(device)
                outputs = model(inputs)
                _, predicted = outputs.max(1)
                total_val += labels.size(0)
                correct_val += (predicted == labels).sum().item()

        val_accuracy = 100 * correct_val / total_val
        val_acc.append(val_accuracy)

        # Print epoch results
        print(f"Epoch [{epoch + 1}/{num_epochs}], Loss: {epoch_loss:.4f}, Train

    return train_acc, val_acc

```

Basic CNN

```
In [39]: class BasicCNN(nn.Module):
    def __init__(self):
        super(BasicCNN, self).__init__()

        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, padding=1)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
        self.pool = nn.MaxPool2d(2, 2)
        self.fc1 = nn.Linear(64 * 8 * 8, 128)
        self.fc2 = nn.Linear(128, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 64 * 8 * 8)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return x
```


```
In [40]: # Instantiate the model, define the Loss function and optimizer
model = BasicCNN().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=0.001, momentum=0.9)

# Train the model and track accuracy
train_acc, val_acc = train_model_with_validation(model, criterion, optimizer, nu

# Plot the training and validation accuracy
plot_training_and_validation_accuracy(train_acc, val_acc, "Basic CNN")

# Evaluate the model on the test set
test_accuracy = evaluate_model(model, testloader)
print(f"Test Accuracy for Basic CNN: {test_accuracy:.2f}%")
```

```
Epoch [1/10], Loss: 2.1246, Train Accuracy: 23.98%, Validation Accuracy: 29.95%
Epoch [2/10], Loss: 1.8434, Train Accuracy: 34.16%, Validation Accuracy: 37.05%
Epoch [3/10], Loss: 1.6887, Train Accuracy: 39.05%, Validation Accuracy: 41.50%
Epoch [4/10], Loss: 1.5935, Train Accuracy: 42.39%, Validation Accuracy: 43.04%
Epoch [5/10], Loss: 1.5270, Train Accuracy: 45.09%, Validation Accuracy: 45.70%
Epoch [6/10], Loss: 1.4768, Train Accuracy: 46.78%, Validation Accuracy: 47.55%
Epoch [7/10], Loss: 1.4347, Train Accuracy: 48.25%, Validation Accuracy: 48.99%
Epoch [8/10], Loss: 1.3982, Train Accuracy: 49.95%, Validation Accuracy: 50.24%
Epoch [9/10], Loss: 1.3604, Train Accuracy: 51.36%, Validation Accuracy: 51.75%
Epoch [10/10], Loss: 1.3310, Train Accuracy: 52.39%, Validation Accuracy: 52.67%
```



Test Accuracy: 55.45%

Test Accuracy for Basic CNN: 55.45%

Hyperparameter Tuning for Basic CNN

```
In [41]: class OptimizedCNN(nn.Module):
def __init__(self):
    super(OptimizedCNN, self).__init__()

    # Convolutional Layers with Batch Normalization and Dropout
    self.conv1 = nn.Conv2d(3, 64, kernel_size=3, padding=1)
    self.bn1 = nn.BatchNorm2d(64)

    self.conv2 = nn.Conv2d(64, 128, kernel_size=3, padding=1)
    self.bn2 = nn.BatchNorm2d(128)

    self.conv3 = nn.Conv2d(128, 128, kernel_size=3, padding=1)
    self.bn3 = nn.BatchNorm2d(128)

    # Pooling and Dropout Layers
    self.pool = nn.MaxPool2d(2, 2)
    self.dropout = nn.Dropout(0.3)

    # Fully Connected Layers
    self.fc1 = nn.Linear(128 * 4 * 4, 256)
    self.fc2 = nn.Linear(256, 128)
    self.fc3 = nn.Linear(128, 10)

def forward(self, x):
    # Convolutional Layer 1
    x = self.pool(F.relu(self.bn1(self.conv1(x))))
```

```

# Convolutional Layer 2
x = self.pool(F.relu(self.bn2(self.conv2(x))))

# Convolutional Layer 3
x = self.pool(F.relu(self.bn3(self.conv3(x))))

# Flatten for Fully Connected Layers
x = x.view(-1, 128 * 4 * 4)

# Fully Connected Layers with Dropout
x = F.relu(self.fc1(x))
x = self.dropout(x)

x = F.relu(self.fc2(x))
x = self.dropout(x)

# Output Layer
x = self.fc3(x)

return x

```

```

In [42]: # Instantiate the model, define the loss function and optimizer
model_optimized = OptimizedCNN().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model_optimized.parameters(), lr=0.01, momentum=0.9,
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='max', pa

```

E:\Codes_data\try\envs\ai\lib\site-packages\torch\optim\lr_scheduler.py:62: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.

warnings.warn(

```

In [44]: # Train the model and track accuracy
train_acc_optimized, val_acc_optimized = train_model_with_validation(
    model_optimized, criterion, optimizer, num_epochs=10
)

# Plot the training and validation accuracy for the optimized model
plot_training_and_validation_accuracy(train_acc_optimized, val_acc_optimized, "C


# Evaluate the optimized model on the test set
test_accuracy_optimized = evaluate_model(model_optimized, testloader)
print(f"Test Accuracy for Optimized CNN: {test_accuracy_optimized:.2f}%")

```

```

Epoch [1/10], Loss: 1.6762, Train Accuracy: 37.65%, Validation Accuracy: 50.49%
Epoch [2/10], Loss: 1.3045, Train Accuracy: 52.37%, Validation Accuracy: 57.79%
Epoch [3/10], Loss: 1.1461, Train Accuracy: 58.85%, Validation Accuracy: 62.52%
Epoch [4/10], Loss: 1.0446, Train Accuracy: 62.90%, Validation Accuracy: 62.61%
Epoch [5/10], Loss: 0.9651, Train Accuracy: 66.41%, Validation Accuracy: 63.07%
Epoch [6/10], Loss: 0.9112, Train Accuracy: 68.00%, Validation Accuracy: 71.30%
Epoch [7/10], Loss: 0.8651, Train Accuracy: 69.45%, Validation Accuracy: 71.00%
Epoch [8/10], Loss: 0.8311, Train Accuracy: 71.05%, Validation Accuracy: 73.83%
Epoch [9/10], Loss: 0.7943, Train Accuracy: 72.39%, Validation Accuracy: 72.42%
Epoch [10/10], Loss: 0.7704, Train Accuracy: 73.32%, Validation Accuracy: 74.12%

```



Test Accuracy: 75.94%

Test Accuracy for Optimized CNN: 75.94%

```
In [50]: # Function to define and summarise ResNet-18
def resnet18_summary():
    print("\nResNet-18 Model Summary:")
    model = models.resnet18(weights=None)
    model.fc = nn.Linear(model.fc.in_features, 10)
    model = model.to(device)
    summary(model, (3, 32, 32))

# Function to define and summarise AlexNet
def alexnet_summary():
    print("\nAlexNet Model Summary:")
    model = models.alexnet(weights=None)
    model.features[0] = nn.Conv2d(3, 64, kernel_size=3, stride=1, padding=1)
    model.classifier[6] = nn.Linear(model.classifier[6].in_features, 10)
    model = model.to(device)
    summary(model, (3, 32, 32))

# Function to define and summarise MobileNetV2
def mobilenet_v2_summary():
    print("\nMobileNetV2 Model Summary:")
    model = models.mobilenet_v2(weights=None)
    model.classifier[1] = nn.Linear(model.classifier[1].in_features, 10)
    model = model.to(device)
    summary(model, (3, 32, 32))
```

Define ResNet-18 & MobileNetV2 Model


```
In [31]: def get_resnet18():
          model = torchvision.models.resnet18(weights=None)
          num_features = model.fc.in_features
          model.fc = nn.Linear(model.fc.in_features, 10)

          model.fc = nn.Sequential(
              nn.Dropout(0.5), # Add dropout with 50% probability
              nn.Linear(num_features, 10)
          )

          return model.to(device)

def get_mobilenet_v2():
    model = torchvision.models.mobilenet_v2(weights=None)
    model.classifier[1] = nn.Linear(model.classifier[1].in_features, 10)
    return model.to(device)

In [45]: # Display resnet18 model summaries
          resnet18_summary()
```

ResNet-18 Model Summary:

Layer (type)	Output Shape	Param #
=====		
Conv2d-1	[-1, 64, 16, 16]	9,408
BatchNorm2d-2	[-1, 64, 16, 16]	128
ReLU-3	[-1, 64, 16, 16]	0
MaxPool2d-4	[-1, 64, 8, 8]	0
Conv2d-5	[-1, 64, 8, 8]	36,864
BatchNorm2d-6	[-1, 64, 8, 8]	128
ReLU-7	[-1, 64, 8, 8]	0
Conv2d-8	[-1, 64, 8, 8]	36,864
BatchNorm2d-9	[-1, 64, 8, 8]	128
ReLU-10	[-1, 64, 8, 8]	0
BasicBlock-11	[-1, 64, 8, 8]	0
Conv2d-12	[-1, 64, 8, 8]	36,864
BatchNorm2d-13	[-1, 64, 8, 8]	128
ReLU-14	[-1, 64, 8, 8]	0
Conv2d-15	[-1, 64, 8, 8]	36,864
BatchNorm2d-16	[-1, 64, 8, 8]	128
ReLU-17	[-1, 64, 8, 8]	0
BasicBlock-18	[-1, 64, 8, 8]	0
Conv2d-19	[-1, 128, 4, 4]	73,728
BatchNorm2d-20	[-1, 128, 4, 4]	256
ReLU-21	[-1, 128, 4, 4]	0
Conv2d-22	[-1, 128, 4, 4]	147,456
BatchNorm2d-23	[-1, 128, 4, 4]	256
Conv2d-24	[-1, 128, 4, 4]	8,192
BatchNorm2d-25	[-1, 128, 4, 4]	256
ReLU-26	[-1, 128, 4, 4]	0
BasicBlock-27	[-1, 128, 4, 4]	0
Conv2d-28	[-1, 128, 4, 4]	147,456
BatchNorm2d-29	[-1, 128, 4, 4]	256
ReLU-30	[-1, 128, 4, 4]	0
Conv2d-31	[-1, 128, 4, 4]	147,456
BatchNorm2d-32	[-1, 128, 4, 4]	256
ReLU-33	[-1, 128, 4, 4]	0
BasicBlock-34	[-1, 128, 4, 4]	0
Conv2d-35	[-1, 256, 2, 2]	294,912
BatchNorm2d-36	[-1, 256, 2, 2]	512
ReLU-37	[-1, 256, 2, 2]	0
Conv2d-38	[-1, 256, 2, 2]	589,824
BatchNorm2d-39	[-1, 256, 2, 2]	512
Conv2d-40	[-1, 256, 2, 2]	32,768
BatchNorm2d-41	[-1, 256, 2, 2]	512
ReLU-42	[-1, 256, 2, 2]	0
BasicBlock-43	[-1, 256, 2, 2]	0
Conv2d-44	[-1, 256, 2, 2]	589,824
BatchNorm2d-45	[-1, 256, 2, 2]	512
ReLU-46	[-1, 256, 2, 2]	0
Conv2d-47	[-1, 256, 2, 2]	589,824
BatchNorm2d-48	[-1, 256, 2, 2]	512
ReLU-49	[-1, 256, 2, 2]	0
BasicBlock-50	[-1, 256, 2, 2]	0
Conv2d-51	[-1, 512, 1, 1]	1,179,648
BatchNorm2d-52	[-1, 512, 1, 1]	1,024
ReLU-53	[-1, 512, 1, 1]	0
Conv2d-54	[-1, 512, 1, 1]	2,359,296
BatchNorm2d-55	[-1, 512, 1, 1]	1,024
Conv2d-56	[-1, 512, 1, 1]	131,072

BatchNorm2d-57	[-1, 512, 1, 1]	1,024
ReLU-58	[-1, 512, 1, 1]	0
BasicBlock-59	[-1, 512, 1, 1]	0
Conv2d-60	[-1, 512, 1, 1]	2,359,296
BatchNorm2d-61	[-1, 512, 1, 1]	1,024
ReLU-62	[-1, 512, 1, 1]	0
Conv2d-63	[-1, 512, 1, 1]	2,359,296
BatchNorm2d-64	[-1, 512, 1, 1]	1,024
ReLU-65	[-1, 512, 1, 1]	0
BasicBlock-66	[-1, 512, 1, 1]	0
AdaptiveAvgPool2d-67	[-1, 512, 1, 1]	0
Linear-68	[-1, 10]	5,130

=====


Total params: 11,181,642
 Trainable params: 11,181,642
 Non-trainable params: 0

Input size (MB): 0.01
 Forward/backward pass size (MB): 1.29
 Params size (MB): 42.65
 Estimated Total Size (MB): 43.95

Baseline Training for ResNet-18

```
In [32]: # Baseline Training for ResNet-18
print("Training ResNet-18 Baseline...")
resnet18_baseline = get_resnet18()
criterion = nn.CrossEntropyLoss()
optimizer_resnet18 = optim.SGD(resnet18_baseline.parameters(), lr=0.001, momentum=0.9)
train_acc_resnet18, val_acc_resnet18 = train_model_with_validation(resnet18_baseline, criterion, optimizer_resnet18, testloader)
plot_training_and_validation_accuracy(train_acc_resnet18, val_acc_resnet18, "ResNet-18 Baseline")
print("Evaluating ResNet-18 Baseline Model on Test Set...")
evaluate_model(resnet18_baseline, testloader)
```

```
Training ResNet-18 Baseline...
Epoch [1/10], Loss: 2.0058, Train Accuracy: 27.62%, Validation Accuracy: 38.71%
Epoch [2/10], Loss: 1.6376, Train Accuracy: 39.27%, Validation Accuracy: 43.98%
Epoch [3/10], Loss: 1.5117, Train Accuracy: 44.18%, Validation Accuracy: 46.63%
Epoch [4/10], Loss: 1.4286, Train Accuracy: 47.62%, Validation Accuracy: 50.05%
Epoch [5/10], Loss: 1.3560, Train Accuracy: 50.60%, Validation Accuracy: 52.27%
Epoch [6/10], Loss: 1.3019, Train Accuracy: 52.69%, Validation Accuracy: 55.09%
Epoch [7/10], Loss: 1.2435, Train Accuracy: 55.14%, Validation Accuracy: 56.06%
Epoch [8/10], Loss: 1.2073, Train Accuracy: 56.63%, Validation Accuracy: 58.05%
Epoch [9/10], Loss: 1.1672, Train Accuracy: 57.99%, Validation Accuracy: 58.34%
Epoch [10/10], Loss: 1.1338, Train Accuracy: 59.37%, Validation Accuracy: 59.44%
```



Evaluating ResNet-18 Baseline Model on Test Set...

Test Accuracy: 61.84%


Out[32]: 61.84

Baseline Training for MobileNetV2

```
In [33]: print("\nTraining MobileNetV2 Baseline...")
mobilenet_v2_baseline = get_mobilenet_v2()
optimizer_mobilenet_v2 = optim.SGD(mobilenet_v2_baseline.parameters(), lr=0.001,
train_acc_mobilenet_v2, val_acc_mobilenet_v2 = train_model_with_validation(mobil
plot_training_and_validation_accuracy(train_acc_mobilenet_v2, val_acc_mobilenet_
print("Evaluating MobileNetV2 Baseline Model on Test Set...")
evaluate_model(mobilenet_v2_baseline, testloader)
```

Training MobileNetV2 Baseline...

Epoch [1/10], Loss: 2.2174, Train Accuracy: 15.89%, Validation Accuracy: 21.15%
 Epoch [2/10], Loss: 1.9694, Train Accuracy: 24.87%, Validation Accuracy: 29.33%
 Epoch [3/10], Loss: 1.8321, Train Accuracy: 31.12%, Validation Accuracy: 33.32%
 Epoch [4/10], Loss: 1.7515, Train Accuracy: 34.29%, Validation Accuracy: 36.15%
 Epoch [5/10], Loss: 1.6960, Train Accuracy: 36.70%, Validation Accuracy: 38.07%
 Epoch [6/10], Loss: 1.6461, Train Accuracy: 38.77%, Validation Accuracy: 39.36%
 Epoch [7/10], Loss: 1.6018, Train Accuracy: 40.38%, Validation Accuracy: 42.00%
 Epoch [8/10], Loss: 1.5622, Train Accuracy: 42.24%, Validation Accuracy: 43.23%
 Epoch [9/10], Loss: 1.5325, Train Accuracy: 43.09%, Validation Accuracy: 43.83%
 Epoch [10/10], Loss: 1.5072, Train Accuracy: 44.47%, Validation Accuracy: 45.15%



Evaluating MobileNetV2 Baseline Model on Test Set...

Test Accuracy: 46.82%

Out[33]: 46.82

```
In [46]: # Model's summary  
mobilenet_v2_summary()
```

MobileNetV2 Model Summary:

Layer (type)	Output Shape	Param #
=====		
Conv2d-1	[-1, 32, 16, 16]	864
BatchNorm2d-2	[-1, 32, 16, 16]	64
ReLU6-3	[-1, 32, 16, 16]	0
Conv2d-4	[-1, 32, 16, 16]	288
BatchNorm2d-5	[-1, 32, 16, 16]	64
ReLU6-6	[-1, 32, 16, 16]	0
Conv2d-7	[-1, 16, 16, 16]	512
BatchNorm2d-8	[-1, 16, 16, 16]	32
InvertedResidual-9	[-1, 16, 16, 16]	0
Conv2d-10	[-1, 96, 16, 16]	1,536
BatchNorm2d-11	[-1, 96, 16, 16]	192
ReLU6-12	[-1, 96, 16, 16]	0
Conv2d-13	[-1, 96, 8, 8]	864
BatchNorm2d-14	[-1, 96, 8, 8]	192
ReLU6-15	[-1, 96, 8, 8]	0
Conv2d-16	[-1, 24, 8, 8]	2,304
BatchNorm2d-17	[-1, 24, 8, 8]	48
InvertedResidual-18	[-1, 24, 8, 8]	0
Conv2d-19	[-1, 144, 8, 8]	3,456
BatchNorm2d-20	[-1, 144, 8, 8]	288
ReLU6-21	[-1, 144, 8, 8]	0
Conv2d-22	[-1, 144, 8, 8]	1,296
BatchNorm2d-23	[-1, 144, 8, 8]	288
ReLU6-24	[-1, 144, 8, 8]	0
Conv2d-25	[-1, 24, 8, 8]	3,456
BatchNorm2d-26	[-1, 24, 8, 8]	48
InvertedResidual-27	[-1, 24, 8, 8]	0
Conv2d-28	[-1, 144, 8, 8]	3,456
BatchNorm2d-29	[-1, 144, 8, 8]	288
ReLU6-30	[-1, 144, 8, 8]	0
Conv2d-31	[-1, 144, 4, 4]	1,296
BatchNorm2d-32	[-1, 144, 4, 4]	288
ReLU6-33	[-1, 144, 4, 4]	0
Conv2d-34	[-1, 32, 4, 4]	4,608
BatchNorm2d-35	[-1, 32, 4, 4]	64
InvertedResidual-36	[-1, 32, 4, 4]	0
Conv2d-37	[-1, 192, 4, 4]	6,144
BatchNorm2d-38	[-1, 192, 4, 4]	384
ReLU6-39	[-1, 192, 4, 4]	0
Conv2d-40	[-1, 192, 4, 4]	1,728
BatchNorm2d-41	[-1, 192, 4, 4]	384
ReLU6-42	[-1, 192, 4, 4]	0
Conv2d-43	[-1, 32, 4, 4]	6,144
BatchNorm2d-44	[-1, 32, 4, 4]	64
InvertedResidual-45	[-1, 32, 4, 4]	0
Conv2d-46	[-1, 192, 4, 4]	6,144
BatchNorm2d-47	[-1, 192, 4, 4]	384
ReLU6-48	[-1, 192, 4, 4]	0
Conv2d-49	[-1, 192, 4, 4]	1,728
BatchNorm2d-50	[-1, 192, 4, 4]	384
ReLU6-51	[-1, 192, 4, 4]	0
Conv2d-52	[-1, 32, 4, 4]	6,144
BatchNorm2d-53	[-1, 32, 4, 4]	64
InvertedResidual-54	[-1, 32, 4, 4]	0
Conv2d-55	[-1, 192, 4, 4]	6,144
BatchNorm2d-56	[-1, 192, 4, 4]	384

ReLU6-57	[-1, 192, 4, 4]	0
Conv2d-58	[-1, 192, 2, 2]	1,728
BatchNorm2d-59	[-1, 192, 2, 2]	384
ReLU6-60	[-1, 192, 2, 2]	0
Conv2d-61	[-1, 64, 2, 2]	12,288
BatchNorm2d-62	[-1, 64, 2, 2]	128
InvertedResidual-63	[-1, 64, 2, 2]	0
Conv2d-64	[-1, 384, 2, 2]	24,576
BatchNorm2d-65	[-1, 384, 2, 2]	768
ReLU6-66	[-1, 384, 2, 2]	0
Conv2d-67	[-1, 384, 2, 2]	3,456
BatchNorm2d-68	[-1, 384, 2, 2]	768
ReLU6-69	[-1, 384, 2, 2]	0
Conv2d-70	[-1, 64, 2, 2]	24,576
BatchNorm2d-71	[-1, 64, 2, 2]	128
InvertedResidual-72	[-1, 64, 2, 2]	0
Conv2d-73	[-1, 384, 2, 2]	24,576
BatchNorm2d-74	[-1, 384, 2, 2]	768
ReLU6-75	[-1, 384, 2, 2]	0
Conv2d-76	[-1, 384, 2, 2]	3,456
BatchNorm2d-77	[-1, 384, 2, 2]	768
ReLU6-78	[-1, 384, 2, 2]	0
Conv2d-79	[-1, 64, 2, 2]	24,576
BatchNorm2d-80	[-1, 64, 2, 2]	128
InvertedResidual-81	[-1, 64, 2, 2]	0
Conv2d-82	[-1, 384, 2, 2]	24,576
BatchNorm2d-83	[-1, 384, 2, 2]	768
ReLU6-84	[-1, 384, 2, 2]	0
Conv2d-85	[-1, 384, 2, 2]	3,456
BatchNorm2d-86	[-1, 384, 2, 2]	768
ReLU6-87	[-1, 384, 2, 2]	0
Conv2d-88	[-1, 64, 2, 2]	24,576
BatchNorm2d-89	[-1, 64, 2, 2]	128
InvertedResidual-90	[-1, 64, 2, 2]	0
Conv2d-91	[-1, 384, 2, 2]	24,576
BatchNorm2d-92	[-1, 384, 2, 2]	768
ReLU6-93	[-1, 384, 2, 2]	0
Conv2d-94	[-1, 384, 2, 2]	3,456
BatchNorm2d-95	[-1, 384, 2, 2]	768
ReLU6-96	[-1, 384, 2, 2]	0
Conv2d-97	[-1, 96, 2, 2]	36,864
BatchNorm2d-98	[-1, 96, 2, 2]	192
InvertedResidual-99	[-1, 96, 2, 2]	0
Conv2d-100	[-1, 576, 2, 2]	55,296
BatchNorm2d-101	[-1, 576, 2, 2]	1,152
ReLU6-102	[-1, 576, 2, 2]	0
Conv2d-103	[-1, 576, 2, 2]	5,184
BatchNorm2d-104	[-1, 576, 2, 2]	1,152
ReLU6-105	[-1, 576, 2, 2]	0
Conv2d-106	[-1, 96, 2, 2]	55,296
BatchNorm2d-107	[-1, 96, 2, 2]	192
InvertedResidual-108	[-1, 96, 2, 2]	0
Conv2d-109	[-1, 576, 2, 2]	55,296
BatchNorm2d-110	[-1, 576, 2, 2]	1,152
ReLU6-111	[-1, 576, 2, 2]	0
Conv2d-112	[-1, 576, 2, 2]	5,184
BatchNorm2d-113	[-1, 576, 2, 2]	1,152
ReLU6-114	[-1, 576, 2, 2]	0
Conv2d-115	[-1, 96, 2, 2]	55,296
BatchNorm2d-116	[-1, 96, 2, 2]	192

InvertedResidual-117	[-1, 96, 2, 2]	0
Conv2d-118	[-1, 576, 2, 2]	55,296
BatchNorm2d-119	[-1, 576, 2, 2]	1,152
ReLU6-120	[-1, 576, 2, 2]	0
Conv2d-121	[-1, 576, 1, 1]	5,184
BatchNorm2d-122	[-1, 576, 1, 1]	1,152
ReLU6-123	[-1, 576, 1, 1]	0
Conv2d-124	[-1, 160, 1, 1]	92,160
BatchNorm2d-125	[-1, 160, 1, 1]	320
InvertedResidual-126	[-1, 160, 1, 1]	0
Conv2d-127	[-1, 960, 1, 1]	153,600
BatchNorm2d-128	[-1, 960, 1, 1]	1,920
ReLU6-129	[-1, 960, 1, 1]	0
Conv2d-130	[-1, 960, 1, 1]	8,640
BatchNorm2d-131	[-1, 960, 1, 1]	1,920
ReLU6-132	[-1, 960, 1, 1]	0
Conv2d-133	[-1, 160, 1, 1]	153,600
BatchNorm2d-134	[-1, 160, 1, 1]	320
InvertedResidual-135	[-1, 160, 1, 1]	0
Conv2d-136	[-1, 960, 1, 1]	153,600
BatchNorm2d-137	[-1, 960, 1, 1]	1,920
ReLU6-138	[-1, 960, 1, 1]	0
Conv2d-139	[-1, 960, 1, 1]	8,640
BatchNorm2d-140	[-1, 960, 1, 1]	1,920
ReLU6-141	[-1, 960, 1, 1]	0
Conv2d-142	[-1, 160, 1, 1]	153,600
BatchNorm2d-143	[-1, 160, 1, 1]	320
InvertedResidual-144	[-1, 160, 1, 1]	0
Conv2d-145	[-1, 960, 1, 1]	153,600
BatchNorm2d-146	[-1, 960, 1, 1]	1,920
ReLU6-147	[-1, 960, 1, 1]	0
Conv2d-148	[-1, 960, 1, 1]	8,640
BatchNorm2d-149	[-1, 960, 1, 1]	1,920
ReLU6-150	[-1, 960, 1, 1]	0
Conv2d-151	[-1, 320, 1, 1]	307,200
BatchNorm2d-152	[-1, 320, 1, 1]	640
InvertedResidual-153	[-1, 320, 1, 1]	0
Conv2d-154	[-1, 1280, 1, 1]	409,600
BatchNorm2d-155	[-1, 1280, 1, 1]	2,560
ReLU6-156	[-1, 1280, 1, 1]	0
Dropout-157	[-1, 1280]	0
Linear-158	[-1, 10]	12,810

=====

Total params: 2,236,682
Trainable params: 2,236,682
Non-trainable params: 0

Input size (MB): 0.01
Forward/backward pass size (MB): 3.13
Params size (MB): 8.53
Estimated Total Size (MB): 11.67

Hyperparameter Tuning for ResNet-18

```
In [34]: # Hyperparameters taken
learning_rate = 0.01
weight_decay = 0.0001
```



```
optimizer_type = 'Adam'

# Function to set optimiser with updated hyperparameters
def get_optimizer(model):
    return optim.SGD(model.parameters(), lr=learning_rate, momentum=0.9, weight_


# Improved training function for ResNet-18 with dropout
def improved_training(model_fn, model_name, num_epochs=10):
    print(f"\nTraining {model_name} with adjusted configuration: lr={learning_ra
    model = model_fn() # Use the modified ResNet-18 with dropout
    optimizer = get_optimizer(model)
    criterion = nn.CrossEntropyLoss()
    train_acc, val_acc = train_model_with_validation(model, criterion, optimizer
    test_accuracy = evaluate_model(model, testloader)
    plot_training_and_validation_accuracy(train_acc, val_acc, f"{model_name}")
    print(f"Test Accuracy for {model_name}: {test_accuracy:.2f}%")
```

```
In [36]: print("Optimisation for ResNet-18 ")
         improved_training(get_resnet18, "ResNet-18")
```

Optimisation for ResNet-18

Training ResNet-18 with adjusted configuration: lr=0.01, optimizer=Adam, weight_decay=0.0005

```
Epoch [1/10], Loss: 1.8988, Train Accuracy: 33.11%, Validation Accuracy: 41.29%
Epoch [2/10], Loss: 1.5775, Train Accuracy: 44.43%, Validation Accuracy: 47.33%
Epoch [3/10], Loss: 1.4215, Train Accuracy: 50.75%, Validation Accuracy: 51.05%
Epoch [4/10], Loss: 1.2880, Train Accuracy: 55.18%, Validation Accuracy: 55.65%
Epoch [5/10], Loss: 1.1726, Train Accuracy: 59.22%, Validation Accuracy: 59.62%
Epoch [6/10], Loss: 1.0730, Train Accuracy: 62.56%, Validation Accuracy: 61.03%
Epoch [7/10], Loss: 1.0177, Train Accuracy: 64.54%, Validation Accuracy: 62.81%
Epoch [8/10], Loss: 0.9590, Train Accuracy: 66.63%, Validation Accuracy: 65.42%
Epoch [9/10], Loss: 0.9172, Train Accuracy: 67.64%, Validation Accuracy: 66.12%
Epoch [10/10], Loss: 0.8687, Train Accuracy: 69.64%, Validation Accuracy: 68.70%
Test Accuracy: 69.71%
```



Test Accuracy for ResNet-18: 69.71%

Hyperparameter Tuning for MobileNetV2

```
In [37]: # Hyperparameters chosen for MobileNetV2
learning_rate = 0.01
weight_decay = 0.0005
optimizer_type = 'Adam'

# Improved training function for MobileNetV2
def improved_training_for_MobileNetV2(model_fn, model_name, num_epochs=10):
    print(f"\nTraining {model_name} with Hyperparameters chosen: lr={learning_rate}, wd={weight_decay}, opt={optimizer_type}")
    model = model_fn()
    optimizer = get_optimizer(model)
    criterion = nn.CrossEntropyLoss()
    train_acc, val_acc = train_model_with_validation(model, criterion, optimizer, num_epochs)
    test_accuracy = evaluate_model(model, testloader)
    plot_training_and_validation_accuracy(train_acc, val_acc, f"{model_name} (Adaptive)")
    print(f"Test Accuracy for {model_name}: {test_accuracy:.2f}%")
```


```
In [38]: print("\n Optimisation for MobileNetV2")
improved_training_for_MobileNetV2(get_mobilenet_v2, "MobileNetV2")
```

Optimisation for MobileNetV2

Training MobileNetV2 with Hyperparameters chosen: lr=0.01, optimizer=Adam, weight_decay=0.0005

Epoch [1/10], Loss: 2.1105, Train Accuracy: 23.05%, Validation Accuracy: 29.43%
Epoch [2/10], Loss: 1.7821, Train Accuracy: 34.42%, Validation Accuracy: 38.31%
Epoch [3/10], Loss: 1.6389, Train Accuracy: 39.89%, Validation Accuracy: 43.57%
Epoch [4/10], Loss: 1.5379, Train Accuracy: 44.20%, Validation Accuracy: 45.73%
Epoch [5/10], Loss: 1.4644, Train Accuracy: 47.15%, Validation Accuracy: 44.17%
Epoch [6/10], Loss: 1.4023, Train Accuracy: 49.05%, Validation Accuracy: 49.78%
Epoch [7/10], Loss: 1.3377, Train Accuracy: 51.70%, Validation Accuracy: 51.77%
Epoch [8/10], Loss: 1.2863, Train Accuracy: 53.87%, Validation Accuracy: 55.05%
Epoch [9/10], Loss: 1.2491, Train Accuracy: 55.12%, Validation Accuracy: 54.91%
Epoch [10/10], Loss: 1.1934, Train Accuracy: 57.34%, Validation Accuracy: 57.78%
Test Accuracy: 60.23%

Training and Validation Accuracy Curve - MobileNetV2 (Adjusted)



Test Accuracy for MobileNetV2: 60.23%

```
In [47]: # Model's summary  
alexnet_summary()
```

AlexNet Model Summary:

Layer (type)	Output Shape	Param #
=====		
Conv2d-1	[-1, 64, 32, 32]	1,792
ReLU-2	[-1, 64, 32, 32]	0
MaxPool2d-3	[-1, 64, 15, 15]	0
Conv2d-4	[-1, 192, 15, 15]	307,392
ReLU-5	[-1, 192, 15, 15]	0
MaxPool2d-6	[-1, 192, 7, 7]	0
Conv2d-7	[-1, 384, 7, 7]	663,936
ReLU-8	[-1, 384, 7, 7]	0
Conv2d-9	[-1, 256, 7, 7]	884,992
ReLU-10	[-1, 256, 7, 7]	0
Conv2d-11	[-1, 256, 7, 7]	590,080
ReLU-12	[-1, 256, 7, 7]	0
MaxPool2d-13	[-1, 256, 3, 3]	0
AdaptiveAvgPool2d-14	[-1, 256, 6, 6]	0
Dropout-15	[-1, 9216]	0
Linear-16	[-1, 4096]	37,752,832
ReLU-17	[-1, 4096]	0
Dropout-18	[-1, 4096]	0
Linear-19	[-1, 4096]	16,781,312
ReLU-20	[-1, 4096]	0
Linear-21	[-1, 10]	40,970
=====		
Total params: 57,023,306		
Trainable params: 57,023,306		
Non-trainable params: 0		

Input size (MB): 0.01		
Forward/backward pass size (MB): 2.83		
Params size (MB): 217.53		
Estimated Total Size (MB): 220.36		

Define AlexNet

```
In [24]: # Define the AlexNet model and move to the device
def get_alexnet():
    model = models.alexnet(pretrained=False)
    model.features[0] = nn.Conv2d(3, 64, kernel_size=3, stride=1, padding=1)
    model.classifier[6] = nn.Linear(model.classifier[6].in_features, 10) # 10 classes
    return model.to(device)

# Training function with accuracy tracking
def train_model_with_accuracy_tracking(model, criterion, optimizer, trainloader,
    train_acc, val_acc = [], []):

    for epoch in range(num_epochs):
        # Training phase
        model.train()
        correct_train, total_train = 0, 0

        for inputs, labels in trainloader:
            inputs, labels = inputs.to(device), labels.to(device)
            optimizer.zero_grad()
            outputs = model(inputs)
```

```

        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        # Track accuracy
        _, predicted = outputs.max(1)
        total_train += labels.size(0)
        correct_train += predicted.eq(labels).sum().item()

    train_accuracy = 100 * correct_train / total_train
    train_acc.append(train_accuracy)

    # Validation phase
    model.eval()
    correct_val, total_val = 0, 0
    with torch.no_grad():
        for inputs, labels in valloader:
            inputs, labels = inputs.to(device), labels.to(device)
            outputs = model(inputs)
            _, predicted = outputs.max(1)
            total_val += labels.size(0)
            correct_val += predicted.eq(labels).sum().item()

    val_accuracy = 100 * correct_val / total_val
    val_acc.append(val_accuracy)

    # Print epoch results
    print(f"Epoch [{epoch+1}/{num_epochs}], Train Accuracy: {train_accuracy:}

return train_acc, val_acc

```

Baseline Training for AlexNet

```

In [25]: # Load model, criterion, and optimizer
model = get_alexnet()
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)

# Assuming trainloader, valloader, and testloader are defined
# Train the model
train_acc, val_acc = train_model_with_accuracy_tracking(model, criterion, optimi


# Plot training and validation accuracy
plot_training_and_validation_accuracy(train_acc, val_acc, "AlexNet Baseline")

# Evaluate model on the test set
print("Evaluating AlexNet Baseline Model on Test Set...")
evaluate_model(model, testloader)

```

Epoch [1/10], Train Accuracy: 10.78%, Validation Accuracy: 9.50%
 Epoch [2/10], Train Accuracy: 10.92%, Validation Accuracy: 13.33%
 Epoch [3/10], Train Accuracy: 16.89%, Validation Accuracy: 22.52%
 Epoch [4/10], Train Accuracy: 23.56%, Validation Accuracy: 24.63%
 Epoch [5/10], Train Accuracy: 26.63%, Validation Accuracy: 28.38%
 Epoch [6/10], Train Accuracy: 30.83%, Validation Accuracy: 34.93%
 Epoch [7/10], Train Accuracy: 34.81%, Validation Accuracy: 35.22%
 Epoch [8/10], Train Accuracy: 37.87%, Validation Accuracy: 39.96%
 Epoch [9/10], Train Accuracy: 41.26%, Validation Accuracy: 44.18%
 Epoch [10/10], Train Accuracy: 44.31%, Validation Accuracy: 46.80%

Training and Validation Accuracy Curve - AlexNet Baseline



Evaluating AlexNet Baseline Model on Test Set...
 Test Accuracy: 48.54%

Out[25]: 48.54

Hyperparameter Tuning for AlexNet

```
In [27]: # Hyperparameter tuning with different learning rate and weight decay
print("\nHyperparameter Tuning: lr=0.01, weight_decay=0.0005")
model_tuned = get_alexnet() # Instantiate a fresh model
optimizer_tuned = optim.SGD(model_tuned.parameters(), lr=0.01, momentum=0.9, wei


# Train the model with tuned hyperparameters
train_acc_tuned, val_acc_tuned = train_model_with_accuracy_tracking(model_tuned,

# Plot training and validation accuracy for tuned model
plot_training_and_validation_accuracy(train_acc_tuned, val_acc_tuned, "AlexNet T

# Evaluate tuned model on the test set
print("Evaluating AlexNet Tuned Model on Test Set...")
evaluate_model(model_tuned, testloader)
```

Hyperparameter Tuning: lr=0.01, weight_decay=0.0005

Epoch [1/10], Train Accuracy: 18.94%, Validation Accuracy: 27.75%
Epoch [2/10], Train Accuracy: 35.12%, Validation Accuracy: 42.53%
Epoch [3/10], Train Accuracy: 46.94%, Validation Accuracy: 49.29%
Epoch [4/10], Train Accuracy: 55.76%, Validation Accuracy: 58.77%
Epoch [5/10], Train Accuracy: 62.10%, Validation Accuracy: 63.97%
Epoch [6/10], Train Accuracy: 66.19%, Validation Accuracy: 67.93%
Epoch [7/10], Train Accuracy: 69.78%, Validation Accuracy: 69.23%
Epoch [8/10], Train Accuracy: 72.31%, Validation Accuracy: 71.42%
Epoch [9/10], Train Accuracy: 75.21%, Validation Accuracy: 72.86%
Epoch [10/10], Train Accuracy: 77.26%, Validation Accuracy: 75.91%



Evaluating AlexNet Tuned Model on Test Set...

Test Accuracy: 77.66%

Out[27]: 77.66

```
In [49]: torch.cuda.empty_cache()  
# print(torch.cuda.memory_summary(device=None, abbreviated=True))
```