

Importing Libraries

```
In [48]: import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, random_split
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt
import itertools
import numpy as np
from torchsummary import summary
from torchvision import models
import torch.nn.functional as F
```

Set Device (GPU)

```
In [2]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
device
```

```
Out[2]: device(type='cuda')
```

Data Transformations for Training and Testing

```
In [3]: # Data Transformations
transform_train = transforms.Compose([
    transforms.RandomCrop(32, padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])

transform_test = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])

mean = [0.4914, 0.4822, 0.4465]
std = [0.2023, 0.1994, 0.2010]
classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse']
```

Loading CIFAR-10 Dataset

```
In [4]: # Load CIFAR-10 dataset
trainset_full = torchvision.datasets.CIFAR10(root='./data', train=True, download=True)
testset = torchvision.datasets.CIFAR10(root='./data', train=False, download=True)

# Display sample images
```

```
def show_images(dataset, classes, mean=mean, std=std):
    fig, axs = plt.subplots(1, 5, figsize=(15, 3))
    for i in range(5):
        img, label = dataset[i]
        img = img * torch.tensor(std).view(3, 1, 1) + torch.tensor(mean).view(3,
        npimg = np.clip(img.numpy(), 0, 1)
        axs[i].imshow(np.transpose(npimg, (1, 2, 0)))
        axs[i].set_title(classes[label])
        axs[i].axis('off')
    plt.show()

show_images(trainset_full, classes)
```

Files already downloaded and verified

Files already downloaded and verified



Split Dataset into Training and Validation Sets (80-20 split)

```
In [5]: # Split into Training and Validation Sets (80-20 split)
train_size = int(0.8 * len(trainset_full))
val_size = len(trainset_full) - train_size
trainset, valset = random_split(trainset_full, [train_size, val_size])

trainloader = DataLoader(trainset, batch_size=128, shuffle=True, num_workers=2)
valloader = DataLoader(valset, batch_size=128, shuffle=False, num_workers=2)
testloader = DataLoader(testset, batch_size=100, shuffle=False, num_workers=2)
```

Dataset Summary and Class Display

```
In [6]: def dataset_summary():
    print("Dataset Summary:")
    print(f"Total training samples: {len(trainset_full)}")
    print(f"Training set size: {len(trainset)}")
    print(f"Validation set size: {len(valset)}")
    print(f"Test set size: {len(testset)}")
    print("Classes:", trainset_full.classes)

dataset_summary()
```

Dataset Summary:

Total training samples: 50000

Training set size: 40000

Validation set size: 10000

Test set size: 10000

Classes: ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

In []:

```
In [46]: class BasicCNN(nn.Module):
    def __init__(self):
        super(BasicCNN, self).__init__()

        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, padding=1)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
        self.pool = nn.MaxPool2d(2, 2)
        self.fc1 = nn.Linear(64 * 8 * 8, 128)
        self.fc2 = nn.Linear(128, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 64 * 8 * 8)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return x
```


```
In [49]: # Instantiate the model, define the Loss function and optimizer
model = BasicCNN().to(device)
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=0.001, momentum=0.9)

# Train the model and track accuracy
train_acc, val_acc = train_model_with_validation(model, criterion, optimizer, num_epochs=10)

# Plot the training and validation accuracy
plot_training_and_validation_accuracy(train_acc, val_acc, "Basic CNN")

# Evaluate the model on the test set
test_accuracy = evaluate_model(model, testloader)
print(f"Test Accuracy for Basic CNN: {test_accuracy:.2f}%")
```

```
Epoch [1/10], Loss: 2.1198, Train Accuracy: 24.10%, Validation Accuracy: 32.41%
Epoch [2/10], Loss: 1.8288, Train Accuracy: 34.57%, Validation Accuracy: 38.93%
Epoch [3/10], Loss: 1.6790, Train Accuracy: 39.48%, Validation Accuracy: 41.33%
Epoch [4/10], Loss: 1.5793, Train Accuracy: 42.66%, Validation Accuracy: 43.80%
Epoch [5/10], Loss: 1.5188, Train Accuracy: 44.86%, Validation Accuracy: 45.98%
Epoch [6/10], Loss: 1.4745, Train Accuracy: 46.64%, Validation Accuracy: 47.60%
Epoch [7/10], Loss: 1.4393, Train Accuracy: 48.23%, Validation Accuracy: 49.63%
Epoch [8/10], Loss: 1.4060, Train Accuracy: 49.41%, Validation Accuracy: 50.13%
Epoch [9/10], Loss: 1.3721, Train Accuracy: 50.66%, Validation Accuracy: 52.45%
Epoch [10/10], Loss: 1.3459, Train Accuracy: 51.83%, Validation Accuracy: 52.17%
```



Test Accuracy: 54.31%

Test Accuracy for Basic CNN: 54.31%

```
In [50]: class OptimizedCNN(nn.Module):
def __init__(self):
    super(OptimizedCNN, self).__init__()

    # Convolutional Layers with Batch Normalization and Dropout
    self.conv1 = nn.Conv2d(3, 64, kernel_size=3, padding=1)
    self.bn1 = nn.BatchNorm2d(64)

    self.conv2 = nn.Conv2d(64, 128, kernel_size=3, padding=1)
    self.bn2 = nn.BatchNorm2d(128)

    self.conv3 = nn.Conv2d(128, 128, kernel_size=3, padding=1)
    self.bn3 = nn.BatchNorm2d(128)

    # Pooling and Dropout Layers
    self.pool = nn.MaxPool2d(2, 2)
    self.dropout = nn.Dropout(0.3)

    # Fully Connected Layers
    self.fc1 = nn.Linear(128 * 4 * 4, 256) # Increased number of neurons
    self.fc2 = nn.Linear(256, 128)
    self.fc3 = nn.Linear(128, 10) # Output Layer for 10 classes

def forward(self, x):
    # Convolutional Layer 1
    x = self.pool(F.relu(self.bn1(self.conv1(x))))

    # Convolutional Layer 2
    x = self.pool(F.relu(self.bn2(self.conv2(x))))
```

```

# Convolutional Layer 3
x = self.pool(F.relu(self.bn3(self.conv3(x))))

# Flatten for Fully Connected Layers
x = x.view(-1, 128 * 4 * 4)

# Fully Connected Layers with Dropout
x = F.relu(self.fc1(x))
x = self.dropout(x)

x = F.relu(self.fc2(x))
x = self.dropout(x)

# Output Layer
x = self.fc3(x)

return x

```

```

In [51]: # Instantiate the model, define the Loss function and optimizer
model_optimized = OptimizedCNN().to(device)
criterion = nn.CrossEntropyLoss()

# Use SGD with momentum and weight decay
optimizer = torch.optim.SGD(model_optimized.parameters(), lr=0.01, momentum=0.9,

# Learning rate scheduler to reduce LR on plateau in validation accuracy
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='max', pa

```

```

In [53]: # Train the model and track accuracy
train_acc_optimized, val_acc_optimized = train_model_with_validation(
    model_optimized, criterion, optimizer, scheduler=scheduler, num_epochs=10
)

# Plot the training and validation accuracy for the optimized model
plot_training_and_validation_accuracy(train_acc_optimized, val_acc_optimized, "C


# Evaluate the optimized model on the test set
test_accuracy_optimized = evaluate_model(model_optimized, testloader)
print(f"Test Accuracy for Optimized CNN: {test_accuracy_optimized:.2f}%")

```

```

Epoch [1/10], Loss: 0.8757, Train Accuracy: 69.51%, Validation Accuracy: 72.27%
Epoch [2/10], Loss: 0.8368, Train Accuracy: 70.81%, Validation Accuracy: 72.53%
Epoch [3/10], Loss: 0.8043, Train Accuracy: 71.86%, Validation Accuracy: 72.64%
Epoch [4/10], Loss: 0.7738, Train Accuracy: 72.99%, Validation Accuracy: 74.03%
Epoch [5/10], Loss: 0.7460, Train Accuracy: 73.91%, Validation Accuracy: 74.94%
Epoch [6/10], Loss: 0.7215, Train Accuracy: 75.28%, Validation Accuracy: 76.29%
Epoch [7/10], Loss: 0.7093, Train Accuracy: 75.49%, Validation Accuracy: 76.68%
Epoch [8/10], Loss: 0.6774, Train Accuracy: 76.55%, Validation Accuracy: 77.48%
Epoch [9/10], Loss: 0.6699, Train Accuracy: 76.94%, Validation Accuracy: 76.43%
Epoch [10/10], Loss: 0.6534, Train Accuracy: 77.54%, Validation Accuracy: 78.27%

```



Test Accuracy: 78.56%

Test Accuracy for Optimized CNN: 78.56%

In []:

In []:

In []:

```
In [7]: # Function to define and summarise ResNet-18
def resnet18_summary():
    print("\nResNet-18 Model Summary:")
    model = models.resnet18(weights=None) # Use ResNet-18 without pretrained weights
    model.fc = nn.Linear(model.fc.in_features, 10) # Adjust for 10 classes (CIFAR-10)
    model = model.to(device)
    summary(model, (3, 32, 32))

# Function to define and summarize AlexNet
def alexnet_summary():
    print("\nAlexNet Model Summary:")
    model = models.alexnet(weights=None)
    model.features[0] = nn.Conv2d(3, 64, kernel_size=3, stride=1, padding=1) # Adjust kernel size
    model.classifier[6] = nn.Linear(model.classifier[6].in_features, 10) # Adjust for 10 classes
    model = model.to(device)
    summary(model, (3, 32, 32))

# Function to define and summarize MobileNetV2
def mobilenet_v2_summary():
    print("\nMobileNetV2 Model Summary:")
    model = models.mobilenet_v2(weights=None)
    model.classifier[1] = nn.Linear(model.classifier[1].in_features, 10) # Adjust for 10 classes
```

```
model = model.to(device)
summary(model, (3, 32, 32))
```

Define ResNet-18 & MobileNetV2 Model

```
In [8]: def get_resnet18():
        model = torchvision.models.resnet18(weights=None)
        num_features = model.fc.in_features
        model.fc = nn.Linear(model.fc.in_features, 10) # Adjust for 10 CIFAR-10 cla

        model.fc = nn.Sequential(
            nn.Dropout(0.5), # Add dropout with 50% probability
            nn.Linear(num_features, 10) # Final layer for 10 classes (CIFAR-10)
        )

        return model.to(device)

def get_mobilenet_v2():
    model = torchvision.models.mobilenet_v2(weights=None)
    model.classifier[1] = nn.Linear(model.classifier[1].in_features, 10)
    return model.to(device)
```

Plot Training and Validation Accuracy Curves

```
In [9]: def plot_training_and_validation_accuracy(train_acc, val_acc, model_name):
        plt.plot(train_acc, label="Training Accuracy")
        plt.plot(val_acc, label="Validation Accuracy")
        plt.xlabel("Epoch")
        plt.ylabel("Accuracy")
        plt.title(f"Training and Validation Accuracy Curve - {model_name}")
        plt.legend()
        plt.show()
```

Evaluate Model on Test Set

```
In [10]: def evaluate_model(model, dataloader):
        model.eval()
        correct, total = 0, 0
        with torch.no_grad():
            for inputs, labels in dataloader:
                inputs, labels = inputs.to(device), labels.to(device)
                outputs = model(inputs)
                _, predicted = outputs.max(1)
                total += labels.size(0)
                correct += (predicted == labels).sum().item()
        accuracy = 100 * correct / total
        print(f"Test Accuracy: {accuracy:.2f}%")
        return accuracy
```

Training Function with Validation

```
In [23]: # Training function with validation
def train_model_with_validation(model, criterion, optimizer, num_epochs=10):
    train_acc, val_acc = [], []
    for epoch in range(num_epochs):
        # Training
        model.train()
        correct_train, total_train, running_loss = 0, 0, 0.0
        for inputs, labels in trainloader:
            inputs, labels = inputs.to(device), labels.to(device)
            optimizer.zero_grad()
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            # Accuracy and Loss tracking
            _, predicted = outputs.max(1)
            total_train += labels.size(0)
            correct_train += (predicted == labels).sum().item()
            running_loss += loss.item()

        # Calculate training accuracy
        epoch_loss = running_loss / len(trainloader)
        epoch_accuracy = 100 * correct_train / total_train
        train_acc.append(epoch_accuracy)

        # Validation
        model.eval()
        correct_val, total_val = 0, 0
        with torch.no_grad():
            for inputs, labels in valloader:
                inputs, labels = inputs.to(device), labels.to(device)
                outputs = model(inputs)
                _, predicted = outputs.max(1)
                total_val += labels.size(0)
                correct_val += (predicted == labels).sum().item()

        val_accuracy = 100 * correct_val / total_val
        val_acc.append(val_accuracy)

        # Print epoch results
        print(f"Epoch [{epoch + 1}/{num_epochs}], Loss: {epoch_loss:.4f}, Train

    return train_acc, val_acc
```

```
In [49]: # Display model summaries
resnet18_summary()
```


ResNet-18 Model Summary:

Layer (type)	Output Shape	Param #
=====		
Conv2d-1	[-1, 64, 16, 16]	9,408
BatchNorm2d-2	[-1, 64, 16, 16]	128
ReLU-3	[-1, 64, 16, 16]	0
MaxPool2d-4	[-1, 64, 8, 8]	0
Conv2d-5	[-1, 64, 8, 8]	36,864
BatchNorm2d-6	[-1, 64, 8, 8]	128
ReLU-7	[-1, 64, 8, 8]	0
Conv2d-8	[-1, 64, 8, 8]	36,864
BatchNorm2d-9	[-1, 64, 8, 8]	128
ReLU-10	[-1, 64, 8, 8]	0
BasicBlock-11	[-1, 64, 8, 8]	0
Conv2d-12	[-1, 64, 8, 8]	36,864
BatchNorm2d-13	[-1, 64, 8, 8]	128
ReLU-14	[-1, 64, 8, 8]	0
Conv2d-15	[-1, 64, 8, 8]	36,864
BatchNorm2d-16	[-1, 64, 8, 8]	128
ReLU-17	[-1, 64, 8, 8]	0
BasicBlock-18	[-1, 64, 8, 8]	0
Conv2d-19	[-1, 128, 4, 4]	73,728
BatchNorm2d-20	[-1, 128, 4, 4]	256
ReLU-21	[-1, 128, 4, 4]	0
Conv2d-22	[-1, 128, 4, 4]	147,456
BatchNorm2d-23	[-1, 128, 4, 4]	256
Conv2d-24	[-1, 128, 4, 4]	8,192
BatchNorm2d-25	[-1, 128, 4, 4]	256
ReLU-26	[-1, 128, 4, 4]	0
BasicBlock-27	[-1, 128, 4, 4]	0
Conv2d-28	[-1, 128, 4, 4]	147,456
BatchNorm2d-29	[-1, 128, 4, 4]	256
ReLU-30	[-1, 128, 4, 4]	0
Conv2d-31	[-1, 128, 4, 4]	147,456
BatchNorm2d-32	[-1, 128, 4, 4]	256
ReLU-33	[-1, 128, 4, 4]	0
BasicBlock-34	[-1, 128, 4, 4]	0
Conv2d-35	[-1, 256, 2, 2]	294,912
BatchNorm2d-36	[-1, 256, 2, 2]	512
ReLU-37	[-1, 256, 2, 2]	0
Conv2d-38	[-1, 256, 2, 2]	589,824
BatchNorm2d-39	[-1, 256, 2, 2]	512
Conv2d-40	[-1, 256, 2, 2]	32,768
BatchNorm2d-41	[-1, 256, 2, 2]	512
ReLU-42	[-1, 256, 2, 2]	0
BasicBlock-43	[-1, 256, 2, 2]	0
Conv2d-44	[-1, 256, 2, 2]	589,824
BatchNorm2d-45	[-1, 256, 2, 2]	512
ReLU-46	[-1, 256, 2, 2]	0
Conv2d-47	[-1, 256, 2, 2]	589,824
BatchNorm2d-48	[-1, 256, 2, 2]	512
ReLU-49	[-1, 256, 2, 2]	0
BasicBlock-50	[-1, 256, 2, 2]	0
Conv2d-51	[-1, 512, 1, 1]	1,179,648
BatchNorm2d-52	[-1, 512, 1, 1]	1,024
ReLU-53	[-1, 512, 1, 1]	0
Conv2d-54	[-1, 512, 1, 1]	2,359,296
BatchNorm2d-55	[-1, 512, 1, 1]	1,024
Conv2d-56	[-1, 512, 1, 1]	131,072

BatchNorm2d-57	[-1, 512, 1, 1]	1,024
ReLU-58	[-1, 512, 1, 1]	0
BasicBlock-59	[-1, 512, 1, 1]	0
Conv2d-60	[-1, 512, 1, 1]	2,359,296
BatchNorm2d-61	[-1, 512, 1, 1]	1,024
ReLU-62	[-1, 512, 1, 1]	0
Conv2d-63	[-1, 512, 1, 1]	2,359,296
BatchNorm2d-64	[-1, 512, 1, 1]	1,024
ReLU-65	[-1, 512, 1, 1]	0
BasicBlock-66	[-1, 512, 1, 1]	0
AdaptiveAvgPool2d-67	[-1, 512, 1, 1]	0
Linear-68	[-1, 10]	5,130

=====


Total params: 11,181,642
 Trainable params: 11,181,642
 Non-trainable params: 0

Input size (MB): 0.01
 Forward/backward pass size (MB): 1.29
 Params size (MB): 42.65
 Estimated Total Size (MB): 43.95

Baseline Training for ResNet-18

```
In [33]: # Baseline Training for ResNet-18
print("Training ResNet-18 Baseline...")
resnet18_baseline = get_resnet18()
criterion = nn.CrossEntropyLoss()
optimizer_resnet18 = optim.SGD(resnet18_baseline.parameters(), lr=0.001, momentum=0.9)
train_acc_resnet18, val_acc_resnet18 = train_model_with_validation(resnet18_baseline, criterion, optimizer_resnet18, testloader)
plot_training_and_validation_accuracy(train_acc_resnet18, val_acc_resnet18, "ResNet-18 Baseline")
print("Evaluating ResNet-18 Baseline Model on Test Set...")
evaluate_model(resnet18_baseline, testloader)
```

```
Training ResNet-18 Baseline...
Epoch [1/10], Loss: 2.0309, Train Accuracy: 27.00%, Validation Accuracy: 38.15%
Epoch [2/10], Loss: 1.6562, Train Accuracy: 38.73%, Validation Accuracy: 43.85%
Epoch [3/10], Loss: 1.5311, Train Accuracy: 43.43%, Validation Accuracy: 47.30%
Epoch [4/10], Loss: 1.4429, Train Accuracy: 47.42%, Validation Accuracy: 50.81%
Epoch [5/10], Loss: 1.3747, Train Accuracy: 50.13%, Validation Accuracy: 52.86%
Epoch [6/10], Loss: 1.3068, Train Accuracy: 52.62%, Validation Accuracy: 55.21%
Epoch [7/10], Loss: 1.2606, Train Accuracy: 54.42%, Validation Accuracy: 55.18%
Epoch [8/10], Loss: 1.2200, Train Accuracy: 56.04%, Validation Accuracy: 56.49%
Epoch [9/10], Loss: 1.1823, Train Accuracy: 57.13%, Validation Accuracy: 58.60%
Epoch [10/10], Loss: 1.1492, Train Accuracy: 58.60%, Validation Accuracy: 60.12%
```



Evaluating ResNet-18 Baseline Model on Test Set...

Test Accuracy: 62.67%


Out[33]: 62.67

Baseline Training for MobileNetV2

```
In [38]: print("\nTraining MobileNetV2 Baseline...")
mobilenet_v2_baseline = get_mobilenet_v2()
optimizer_mobilenet_v2 = optim.SGD(mobilenet_v2_baseline.parameters(), lr=0.001,
train_acc_mobilenet_v2, val_acc_mobilenet_v2 = train_model_with_validation(mobil
plot_training_and_validation_accuracy(train_acc_mobilenet_v2, val_acc_mobilenet_
print("Evaluating MobileNetV2 Baseline Model on Test Set...")
evaluate_model(mobilenet_v2_baseline, testloader)
```

Training MobileNetV2 Baseline...

```
Epoch [1/10], Loss: 2.2210, Train Accuracy: 16.04%, Validation Accuracy: 21.63%
Epoch [2/10], Loss: 1.9493, Train Accuracy: 26.17%, Validation Accuracy: 30.96%
Epoch [3/10], Loss: 1.8167, Train Accuracy: 31.12%, Validation Accuracy: 33.46%
Epoch [4/10], Loss: 1.7334, Train Accuracy: 34.59%, Validation Accuracy: 36.99%
Epoch [5/10], Loss: 1.6726, Train Accuracy: 37.53%, Validation Accuracy: 38.94%
Epoch [6/10], Loss: 1.6254, Train Accuracy: 39.24%, Validation Accuracy: 41.82%
Epoch [7/10], Loss: 1.5826, Train Accuracy: 40.88%, Validation Accuracy: 43.47%
Epoch [8/10], Loss: 1.5489, Train Accuracy: 42.52%, Validation Accuracy: 44.34%
Epoch [9/10], Loss: 1.5196, Train Accuracy: 43.92%, Validation Accuracy: 45.81%
Epoch [10/10], Loss: 1.4856, Train Accuracy: 45.13%, Validation Accuracy: 46.38%
```



Evaluating MobileNetV2 Baseline Model on Test Set...

Test Accuracy: 48.96%

Out[38]: 48.96

```
In [37]: # Model's summary
# mobilenet_v2_summary()
```

Hyperparameter Tuning for ResNet-18

```
In [34]: # Hyperparameters taken
learning_rate = 0.01
weight_decay = 0.0001
optimizer_type = 'Adam'

# Function to set optimiser with updated hyperparameters
def get_optimizer(model):
    return optim.SGD(model.parameters(), lr=learning_rate, momentum=0.9, weight_


# Improved training function for ResNet-18 with dropout
def improved_training(model_fn, model_name, num_epochs=10):
    print(f"\nTraining {model_name} with adjusted configuration: lr={learning_ra
    model = model_fn() # Use the modified ResNet-18 with dropout
    optimizer = get_optimizer(model)
    criterion = nn.CrossEntropyLoss()
    train_acc, val_acc = train_model_with_validation(model, criterion, optimizer
    test_accuracy = evaluate_model(model, testloader)
    plot_training_and_validation_accuracy(train_acc, val_acc, f"{model_name}")
    print(f"Test Accuracy for {model_name}: {test_accuracy:.2f}%")
```

```
In [35]: print("Optimisation for ResNet-18 ")
         improved_training(get_resnet18, "ResNet-18")
```

Optimisation for ResNet-18

Training ResNet-18 with adjusted configuration: lr=0.01, optimizer=Adam, weight_decay=0.0001

Epoch [1/10], Loss: 1.9281, Train Accuracy: 32.25%, Validation Accuracy: 41.40%
 Epoch [2/10], Loss: 1.5858, Train Accuracy: 44.19%, Validation Accuracy: 48.52%
 Epoch [3/10], Loss: 1.4346, Train Accuracy: 50.22%, Validation Accuracy: 54.10%
 Epoch [4/10], Loss: 1.2861, Train Accuracy: 55.24%, Validation Accuracy: 57.55%
 Epoch [5/10], Loss: 1.2065, Train Accuracy: 58.97%, Validation Accuracy: 62.12%
 Epoch [6/10], Loss: 1.1078, Train Accuracy: 61.87%, Validation Accuracy: 62.55%
 Epoch [7/10], Loss: 1.0480, Train Accuracy: 63.49%, Validation Accuracy: 64.93%
 Epoch [8/10], Loss: 0.9821, Train Accuracy: 65.92%, Validation Accuracy: 66.18%
 Epoch [9/10], Loss: 0.9307, Train Accuracy: 67.58%, Validation Accuracy: 66.69%
 Epoch [10/10], Loss: 0.8938, Train Accuracy: 68.71%, Validation Accuracy: 68.89%
 Test Accuracy: 69.67%



Test Accuracy for ResNet-18: 69.67%

Hyperparameter Tuning for MobileNetV2

```
In [40]: # Hyperparameters chosen for MobileNetV2
         learning_rate = 0.01
         weight_decay = 0.0005
         optimizer_type = 'Adam'

         # Improved training function for MobileNetV2
         def improved_training_for_MobileNetV2(model_fn, model_name, num_epochs=10):
```

```
print(f"\nTraining {model_name} with Hyperparameters chosen: lr={learning_rate}")
model = model_fn()
optimizer = get_optimizer(model)
criterion = nn.CrossEntropyLoss()
train_acc, val_acc = train_model_with_validation(model, criterion, optimizer)
test_accuracy = evaluate_model(model, testloader)
plot_training_and_validation_accuracy(train_acc, val_acc, f"{model_name} (Adjusted)")
print(f"Test Accuracy for {model_name}: {test_accuracy:.2f}%")
```


```
In [41]: print("\n Optimisation for MobileNetV2")
improved_training_for_MobileNetV2(get_mobilenet_v2, "MobileNetV2")
```

Optimisation for MobileNetV2

Training MobileNetV2 with Hyperparameters chosen: lr=0.01, optimizer=Adam, weight_decay=0.0005

Epoch [1/10], Loss: 2.0618, Train Accuracy: 23.90%, Validation Accuracy: 31.32%
 Epoch [2/10], Loss: 1.7527, Train Accuracy: 35.19%, Validation Accuracy: 38.30%
 Epoch [3/10], Loss: 1.6020, Train Accuracy: 40.95%, Validation Accuracy: 41.79%
 Epoch [4/10], Loss: 1.4996, Train Accuracy: 45.10%, Validation Accuracy: 47.01%
 Epoch [5/10], Loss: 1.4370, Train Accuracy: 47.61%, Validation Accuracy: 50.75%
 Epoch [6/10], Loss: 1.3823, Train Accuracy: 49.84%, Validation Accuracy: 52.40%
 Epoch [7/10], Loss: 1.3284, Train Accuracy: 52.04%, Validation Accuracy: 53.23%
 Epoch [8/10], Loss: 1.2844, Train Accuracy: 53.61%, Validation Accuracy: 55.13%
 Epoch [9/10], Loss: 1.2350, Train Accuracy: 55.78%, Validation Accuracy: 56.34%
 Epoch [10/10], Loss: 1.1907, Train Accuracy: 57.40%, Validation Accuracy: 57.76%
 Test Accuracy: 59.73%

Training and Validation Accuracy Curve - MobileNetV2 (Adjusted)



Test Accuracy for MobileNetV2: 59.73%

Define AlexNet

```

In [61]: # Function to create AlexNet model for CIFAR-10
def get_alexnet():
    model = models.alexnet(weights=None)

    # Modify the first convolutional layer to handle 32x32 images
    model.features[0] = nn.Conv2d(3, 64, kernel_size=3, stride=1, padding=1)

    # Calculate flattened size after feature extraction layers
    dummy_input = torch.randn(1, 3, 32, 32).to(device)
    features_output = model.features(dummy_input)
    flattened_size = features_output.view(-1).shape[0]

    # Modify fully connected layers to accommodate CIFAR-10's 10 classes
    model.classifier[1] = nn.Linear(flattened_size, 4096)
    model.classifier[6] = nn.Linear(4096, 10) # Output layer for 10 classes

    return model.to(device) # Move model to the appropriate device

In [62]: # Training function with device handling
def train_model_with_loss_tracking(model, criterion, optimizer, num_epochs=10):
    train_losses = []
    val_losses = []

    for epoch in range(num_epochs):
        # Training phase
        model.train()
        running_loss = 0.0
        for inputs, labels in trainloader:
            # Move inputs and labels to the correct device
            inputs, labels = inputs.to(device), labels.to(device)

            optimizer.zero_grad()
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            running_loss += loss.item() * inputs.size(0)

        epoch_train_loss = running_loss / len(trainloader.dataset)
        train_losses.append(epoch_train_loss)

        # Validation phase
        model.eval()
        running_loss = 0.0
        with torch.no_grad():
            for inputs, labels in valloader:
                # Move inputs and labels to the correct device
                inputs, labels = inputs.to(device), labels.to(device)

                outputs = model(inputs)
                loss = criterion(outputs, labels)
                running_loss += loss.item() * inputs.size(0)

        epoch_val_loss = running_loss / len(valloader.dataset)
        val_losses.append(epoch_val_loss)

    print(f"Epoch [{epoch+1}/{num_epochs}], Train Loss: {epoch_train_loss:.4

```

```
return train_losses, val_losses
```

```
In [39]: # Model's summary  
# alexnet_summary()
```

Baseline Training for AlexNet

```
In [63]: # Initialize and train AlexNet model  
print("Training AlexNet...")  
alexnet_model = get_alexnet() # Move AlexNet to the appropriate device  
criterion = nn.CrossEntropyLoss()  
optimizer = optim.SGD(alexnet_model.parameters(), lr=0.01, momentum=0.9, weight_  
  
# Train the model and get loss values for learning curve  
train_losses, val_losses = train_model_with_loss_tracking(alexnet_model, criteri  
  
# Plot the learning curve  
plot_learning_curve(train_losses, val_losses)
```

Training AlexNet...


```

-----
RuntimeError                                Traceback (most recent call last)
Cell In[63], line 3
      1 # Initialize and train AlexNet model
      2 print("Training AlexNet...")
----> 3 alexnet_model = get_alexnet() # Move AlexNet to the appropriate device
      4 criterion = nn.CrossEntropyLoss()
      5 optimizer = optim.SGD(alexnet_model.parameters(), lr=0.01, momentum=0.9,
weight_decay=5e-4)

Cell In[61], line 10, in get_alexnet()
      8 # Calculate flattened size after feature extraction layers
      9 dummy_input = torch.randn(1, 3, 32, 32).to(device)
--> 10 features_output = model.features(dummy_input)
     11 flattened_size = features_output.view(-1).shape[0]
     12 # Modify fully connected layers to accommodate CIFAR-10's 10 classes

File E:\Codes_data\try\envs\ai\lib\site-packages\torch\nn\modules\module.py:1736,
in Module._wrapped_call_impl(self, *args, **kwargs)
     1734     return self._compiled_call_impl(*args, **kwargs) # type: ignore[mis
c]
     1735 else:
-> 1736     return self._call_impl(*args, **kwargs)

File E:\Codes_data\try\envs\ai\lib\site-packages\torch\nn\modules\module.py:1747,
in Module._call_impl(self, *args, **kwargs)
     1742 # If we don't have any hooks, we want to skip the rest of the logic in
     1743 # this function, and just call forward.
     1744 if not (self._backward_hooks or self._backward_pre_hooks or self._forward
_hooks or self._forward_pre_hooks
     1745         or _global_backward_pre_hooks or _global_backward_hooks
     1746         or _global_forward_hooks or _global_forward_pre_hooks):
-> 1747     return forward_call(*args, **kwargs)
     1749 result = None
     1750 called_always_called_hooks = set()

File E:\Codes_data\try\envs\ai\lib\site-packages\torch\nn\modules\container.py:25
0, in Sequential.forward(self, input)
     248 def forward(self, input):
     249     for module in self:
--> 250         input = module(input)
     251     return input

File E:\Codes_data\try\envs\ai\lib\site-packages\torch\nn\modules\module.py:1736,
in Module._wrapped_call_impl(self, *args, **kwargs)
     1734     return self._compiled_call_impl(*args, **kwargs) # type: ignore[mis
c]
     1735 else:
-> 1736     return self._call_impl(*args, **kwargs)

File E:\Codes_data\try\envs\ai\lib\site-packages\torch\nn\modules\module.py:1747,
in Module._call_impl(self, *args, **kwargs)
     1742 # If we don't have any hooks, we want to skip the rest of the logic in
     1743 # this function, and just call forward.
     1744 if not (self._backward_hooks or self._backward_pre_hooks or self._forward
_hooks or self._forward_pre_hooks
     1745         or _global_backward_pre_hooks or _global_backward_hooks
     1746         or _global_forward_hooks or _global_forward_pre_hooks):
-> 1747     return forward_call(*args, **kwargs)
     1749 result = None

```

```

1750 called_always_called_hooks = set()

File E:\Codes_data\try\envs\ai\lib\site-packages\torch\nn\modules\conv.py:554, in
Conv2d.forward(self, input)
    553 def forward(self, input: Tensor) -> Tensor:
--> 554     return self._conv_forward(input, self.weight, self.bias)

File E:\Codes_data\try\envs\ai\lib\site-packages\torch\nn\modules\conv.py:549, in
Conv2d._conv_forward(self, input, weight, bias)
    537 if self.padding_mode != "zeros":
    538     return F.conv2d(
    539         F.pad(
    540             input, self._reversed_padding_repeated_twice, mode=self.paddi
ng_mode
    (...
    547         self.groups,
    548     )
--> 549 return F.conv2d(
    550     input, weight, bias, self.stride, self.padding, self.dilation, self.g
roups
    551 )

RuntimeError: Input type (torch.cuda.FloatTensor) and weight type (torch.FloatTen
sor) should be the same

```

Hyperparameter Tuning for AlexNet

```

In [19]: print("Hyperparameter Tuning for AlexNet")
alexnet_tuned = get_alexnet()
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(alexnet_tuned.parameters(), lr=0.001, weight_decay=1e-4)

# Learning rate scheduler - Reduce on plateau
scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='max', patience

# Train and validate with improved configuration
train_acc, val_acc = train_model_with_validation(
    model=alexnet_tuned,
    criterion=criterion,
    optimizer=optimizer,
    num_epochs=10,
    scheduler=scheduler
)
plot_training_and_validation_accuracy(train_acc, val_acc, "AlexNet (Improved wit
print("Evaluating AlexNet with Hyperparameter Tuning on Test Set...")
evaluate_model(alexnet_tuned, testloader)

```

Hyperparameter Tuning for AlexNet

```

E:\Codes_data\try\envs\ai\lib\site-packages\torch\optim\lr_scheduler.py:62: UserW
arning: The verbose parameter is deprecated. Please use get_last_lr() to access t
he learning rate.
warnings.warn(

```

```

-----
RuntimeError                                Traceback (most recent call last)
Cell In[19], line 10
      7 scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='max', p
atience=2, factor=0.5, verbose=True)
      9 # Train and validate with improved configuration
--> 10 train_acc, val_acc = train_model_with_validation(
    11     model=alexnet_tuned,
    12     criterion=criterion,
    13     optimizer=optimizer,
    14     num_epochs=10,
    15     scheduler=scheduler
    16 )
    17 plot_training_and_validation_accuracy(train_acc, val_acc, "AlexNet (Impro
ved with Scheduler)")
    18 print("Evaluating AlexNet with Hyperparameter Tuning on Test Set...")

Cell In[12], line 11, in train_model_with_validation(model, criterion, optimizer,
num_epochs, scheduler)
      9 inputs, labels = inputs.to(device), labels.to(device)
     10 optimizer.zero_grad()
--> 11 outputs = model(inputs)
     12 loss = criterion(outputs, labels)
     13 loss.backward()

File E:\Codes_data\try\envs\ai\lib\site-packages\torch\nn\modules\module.py:1736,
in Module._wrapped_call_impl(self, *args, **kwargs)
    1734     return self._compiled_call_impl(*args, **kwargs) # type: ignore[mis
c]
    1735 else:
-> 1736     return self._call_impl(*args, **kwargs)

File E:\Codes_data\try\envs\ai\lib\site-packages\torch\nn\modules\module.py:1747,
in Module._call_impl(self, *args, **kwargs)
    1742 # If we don't have any hooks, we want to skip the rest of the logic in
    1743 # this function, and just call forward.
    1744 if not (self._backward_hooks or self._backward_pre_hooks or self._forward
_hooks or self._forward_pre_hooks
    1745         or _global_backward_pre_hooks or _global_backward_hooks
    1746         or _global_forward_hooks or _global_forward_pre_hooks):
-> 1747     return forward_call(*args, **kwargs)
    1749 result = None
    1750 called_always_called_hooks = set()

File E:\Codes_data\try\envs\ai\lib\site-packages\torchvision\models\alexnet.py:5
1, in AlexNet.forward(self, x)
     49 x = self.avgpool(x)
     50 x = torch.flatten(x, 1)
--> 51 x = self.classifier(x)
     52 return x

File E:\Codes_data\try\envs\ai\lib\site-packages\torch\nn\modules\module.py:1736,
in Module._wrapped_call_impl(self, *args, **kwargs)
    1734     return self._compiled_call_impl(*args, **kwargs) # type: ignore[mis
c]
    1735 else:
-> 1736     return self._call_impl(*args, **kwargs)

File E:\Codes_data\try\envs\ai\lib\site-packages\torch\nn\modules\module.py:1747,
in Module._call_impl(self, *args, **kwargs)

```

```

1742 # If we don't have any hooks, we want to skip the rest of the logic in
1743 # this function, and just call forward.
1744 if not (self._backward_hooks or self._backward_pre_hooks or self._forward
_hooks or self._forward_pre_hooks
1745         or _global_backward_pre_hooks or _global_backward_hooks
1746         or _global_forward_hooks or _global_forward_pre_hooks):
-> 1747     return forward_call(*args, **kwargs)
1749 result = None
1750 called_always_called_hooks = set()

```

File E:\Codes_data\try\envs\ai\lib\site-packages\torch\nn\modules\container.py:250, in Sequential.forward(self, input)

```

248 def forward(self, input):
249     for module in self:
-> 250         input = module(input)
251     return input

```

File E:\Codes_data\try\envs\ai\lib\site-packages\torch\nn\modules\module.py:1736, in Module._wrapped_call_impl(self, *args, **kwargs)

```

1734     return self._compiled_call_impl(*args, **kwargs) # type: ignore[misc]
1735 else:
-> 1736     return self._call_impl(*args, **kwargs)

```

File E:\Codes_data\try\envs\ai\lib\site-packages\torch\nn\modules\module.py:1747, in Module._call_impl(self, *args, **kwargs)

```

1742 # If we don't have any hooks, we want to skip the rest of the logic in
1743 # this function, and just call forward.
1744 if not (self._backward_hooks or self._backward_pre_hooks or self._forward
_hooks or self._forward_pre_hooks
1745         or _global_backward_pre_hooks or _global_backward_hooks
1746         or _global_forward_hooks or _global_forward_pre_hooks):
-> 1747     return forward_call(*args, **kwargs)
1749 result = None
1750 called_always_called_hooks = set()

```

File E:\Codes_data\try\envs\ai\lib\site-packages\torch\nn\modules\linear.py:125, in Linear.forward(self, input)

```

124 def forward(self, input: Tensor) -> Tensor:
-> 125     return F.linear(input, self.weight, self.bias)

```

RuntimeError: mat1 and mat2 shapes cannot be multiplied (128x9216 and 4096x4096)

In []:

In []:

```

In [18]: torch.cuda.empty_cache()
print(torch.cuda.memory_summary(device=None, abbreviated=True))

```

PyTorch CUDA memory summary, device ID 0				
CUDA OOMs: 0		cudaMalloc retries: 0		
Metric	Cur Usage	Peak Usage	Tot Alloc	Tot Freed
Allocated memory	296269 KiB	331421 KiB	799 MiB	522517 KiB
Active memory	296269 KiB	331421 KiB	799 MiB	522517 KiB
Requested memory	296220 KiB	331420 KiB	799 MiB	522515 KiB
GPU reserved memory	352256 KiB	354304 KiB	354304 KiB	2048 KiB
Non-releasable memory	55986 KiB	105315 KiB	409277 KiB	353290 KiB
Allocations	33	38	81	48
Active allocs	33	38	81	48
GPU reserved segments	6	7	7	1
Non-releasable allocs	4	6	16	12
Oversize allocations	0	0	0	0
Oversize GPU segments	0	0	0	0